


## Exercise week 2 - October 6th, 2022 - Exercise 5.

### Commented Solution

Let us consider a distributed system composed of  $N$  processes  $p_1, p_2, \dots, p_n$  each one having a unique integer identifier. Processes are arranged in line topology as in the following figure



Let us assume that there are no failures in the system (i.e., processes are always correct) and that topology links are implemented through perfect point-to-point links.

Write the pseudo-code of a distributed algorithm that is able to build the abstraction of a perfect point-to-point link between any pair of processes (also between those that are not directly connected).

### General Comments

The text says that processes have access to a perfect point-to-point link primitive, let us refer to this component with  $\mathcal{PL}_l$ . This primitive allows every process to exchange messages with at most two other processes, the neighbor processes in the line.

Let us recall the events and properties of a perfect point-to-point primitive.

---

**Module:**

**Name:** PerfectPointToPointLinks, **instance**  $pl$ .

**Events:**

**Request:**  $\langle pl, Send \mid q, m \rangle$ : Requests to send message  $m$  to process  $q$ .

**Indication:**  $\langle pl, Deliver \mid p, m \rangle$ : Delivers message  $m$  sent by process  $p$ .

**Properties:**

**PL1:** *Reliable delivery*: If a correct process  $p$  sends a message  $m$  to a correct process  $q$ , then  $q$  eventually delivers  $m$ .

**PL2:** *No duplication*: No message is delivered by a process more than once.

**PL3:** *No creation*: If some process  $q$  delivers a message  $m$  with sender  $p$ , then  $m$  was previously sent to  $q$  by process  $p$ .

Every process can thus trigger *Send* events and catch *Deliver* events of the  $\mathcal{PL}_l$  component.

The exercise requests to define a protocol a link abstraction, let us name it  $\mathcal{PL}_g$ , that implements a perfect point-to-point link between all pairs of processes. In the initial setting, every process can exchange messages only with at most two processes, namely the neighbor processes on the line. It follows that if two processes are not linked by  $\mathcal{PL}_l$ , such as processes  $p_1$  and  $p_4$ , they cannot exchange messages in the current setting. The target of the exercise is the definition of a protocol that implements a communication primitive guaranteeing all the properties of a perfect point-to-point link between all pairs of processes.

## Protocol Idea

Each process is associated with an integer identifier and it is placed in a topology (the line) such that all the processes with a lower identifier are placed on one side (from the prospective of a selected process) and all processes with a higher identifier are located on the other side. Furthermore, the processes are placed in order on the line with respect to their identifiers. It follows that if process  $p_1$  wants to communicate with process  $p_4$ , it can send its messages to process  $p_2$ , that will relay such messages to  $p_3$ , which finally forwards them to  $p_4$ . More in detail, if the destination of a message is a process with a higher identifier with respect to a selected process, then it is sufficient to forward the message to the neighbors with greater ID till reaching the destination process, and vice-versa for a destination with a lower identifier.

## Pseudo-code

---

### Algorithm 1 $\mathcal{PL}_g$

---

```

1: procedure INIT
2:    $ID \leftarrow \text{unique integer identifier}$ 

3: upon event  $\langle \mathcal{PL}_g, \text{Send} \mid \text{dest}, m \rangle$  do
4:   if  $\text{dest} > ID$  then
5:     trigger  $\langle \mathcal{PL}_l, \text{Send} \mid ID + 1, \langle ID, \text{dest}, m \rangle \rangle$ 
6:   else if  $\text{dest} < ID$  then
7:     trigger  $\langle \mathcal{PL}_l, \text{Send} \mid ID - 1, \langle ID, \text{dest}, m \rangle \rangle$ 
8:   else
9:     trigger  $\langle \mathcal{PL}_l, \text{Send} \mid ID, \langle ID, \text{dest}, m \rangle \rangle$ 

10: upon event  $\langle \mathcal{PL}_l, \text{Deliver} \mid q, \langle \text{source}, \text{dest}, m \rangle \rangle$  do
11:   if  $\text{dest} == ID$  then
12:     trigger  $\langle \mathcal{PL}_g, \text{Deliver} \mid \text{source}, m \rangle$ 
13:   else if  $\text{dest} > ID$  then
14:     trigger  $\langle \mathcal{PL}_l, \text{Send} \mid ID + 1, \langle \text{source}, \text{dest}, m \rangle \rangle$ 
15:   else
16:     trigger  $\langle \mathcal{PL}_l, \text{Send} \mid ID - 1, \langle \text{source}, \text{dest}, m \rangle \rangle$ 

```

---

## Pseudo-code Comments

The first upon block (lines 3-9) rules how a process must act when the *Send* operation of the global perfect point-to-point primitive we are developing,  $\mathcal{PL}_g$ , is triggered. It is a tiny procedure that simply starts the propagation of a new message,  $\langle \text{source}, \text{dest}, m \rangle$ , over the proper link,  $\mathcal{PL}_l$ . Again, a content  $m$  (let use this word to distinguish  $m$  from  $\langle \text{source}, \text{dest}, m \rangle$ ) targeted to a process with an higher identifier ( $\text{dest}$ ) is reachable (on the line) by relaying the content to the neighbor with higher ID and vice-versa.

The second upon block rules how a process that has received a message  $\langle \text{source}, \text{dest}, m \rangle$  from the primitive  $\mathcal{PL}_l$  (namely the local perfect point-to-point link) must act. This procedure compares the  $\text{dest}$  field contained inside the received message with the process identifier. If the process is the destination of a content, then it triggers the *Deliver* operation of the primitive we are developing (content  $m$  was targeted to this process). Otherwise, it continues the propagation of the message  $\langle \text{source}, \text{dest}, m \rangle$  over the proper perfect point-to-point link. Notice that we need to forward the information about the source ID  $\text{source}$  and the destination id  $\text{dest}$  insider the message, otherwise a process (not linked with the source process) receiving a content  $m$  from  $\mathcal{PL}_l$  cannot assert whether the content is targeted to it-self and which process was its source.

## Informal Correctness Proofs

We briefly check the correctness of the provided pseudo-code. The  $\mathcal{PL}_g$  primitive we defined must handle two events,  $\langle \mathcal{PL}_g, \text{Send} \mid \text{dest}, m \rangle$  and  $\langle \mathcal{PL}_g, \text{Deliver} \mid \text{source}, m \rangle$ . Furthermore, it must guarantee the *reliable delivery, no duplication* and *No creation* property of a perfect point-to-point link abstraction.

*Reliable delivery:* if a  $\langle \mathcal{PL}_g, \text{Send} \mid \text{dest}, m \rangle$  event is generated (from the application layer), then the first upon procedure reacts to this event by starting the propagation of the message  $\langle \text{source}, \text{dest}, m \rangle$  over the proper perfect point-to-point link, triggering a *Send* event on the  $\mathcal{PL}_l$  primitive. The  $\mathcal{PL}_l$  abstraction guarantees all the perfect point-to-point link properties between two linked processes. Therefore, every *Send* event of the  $\mathcal{PL}_l$  component generates a *Deliver* event on the destination process (e.g. a  $\langle \mathcal{PL}_l, \text{Send} \mid p_2, \langle \text{source}, \text{dest}, m \rangle \rangle$  event on process  $p_1$  generates the  $\langle \mathcal{PL}_l, \text{Deliver} \mid p_1, \langle \text{source}, \text{dest}, m \rangle \rangle$  event on process  $p_2$ ). The second upon procedure guarantees the propagation of the message  $\langle \text{source}, \text{dest}, m \rangle$  till process  $p_i, i = \text{dest}$ . Only process  $p_i, i = \text{dest}$  triggers the  $\langle \mathcal{PL}_g, \text{Deliver} \mid \text{source}, m \rangle$  event. The propagation policy eventually guarantees the occurrence of this event.

*No duplication:*  $\mathcal{PL}_l$  guarantees no duplication of the messages diffused by the primitive. For a single  $\langle \mathcal{PL}_g, \text{Send} \mid \text{dest}, m \rangle$  event a single message  $\langle \text{source}, \text{dest}, m \rangle$  is generated and propagated over the line using  $\mathcal{PL}_l$ , furthermore every process relays at most once a single  $\langle \text{source}, \text{dest}, m \rangle$  message.

*No creation:*  $\langle \mathcal{PL}_g, \text{Deliver} \mid \text{source}, m \rangle$  event can be triggered only from the reception of a message  $\langle \text{source}, \text{dest}, m \rangle$ . The message  $\langle \text{source}, \text{dest}, m \rangle$  is initially generated by the processes only managing a  $\langle \mathcal{PL}_g, \text{Send} \mid \text{dest}, m \rangle$  event.