

Dependable Distributed Systems

Master of Science in Engineering in Computer Science

AA 2022/2023

LECTURE 6: FAILURE DETECTION

Recap on Timing Assumptions

Synchronous

- timing assumptions are explicit on
 - Upper bounds on process execution time
 - Upper bounds on communication time
 - Existence of a common global clock

Asynchronous

- there are no timing assumptions

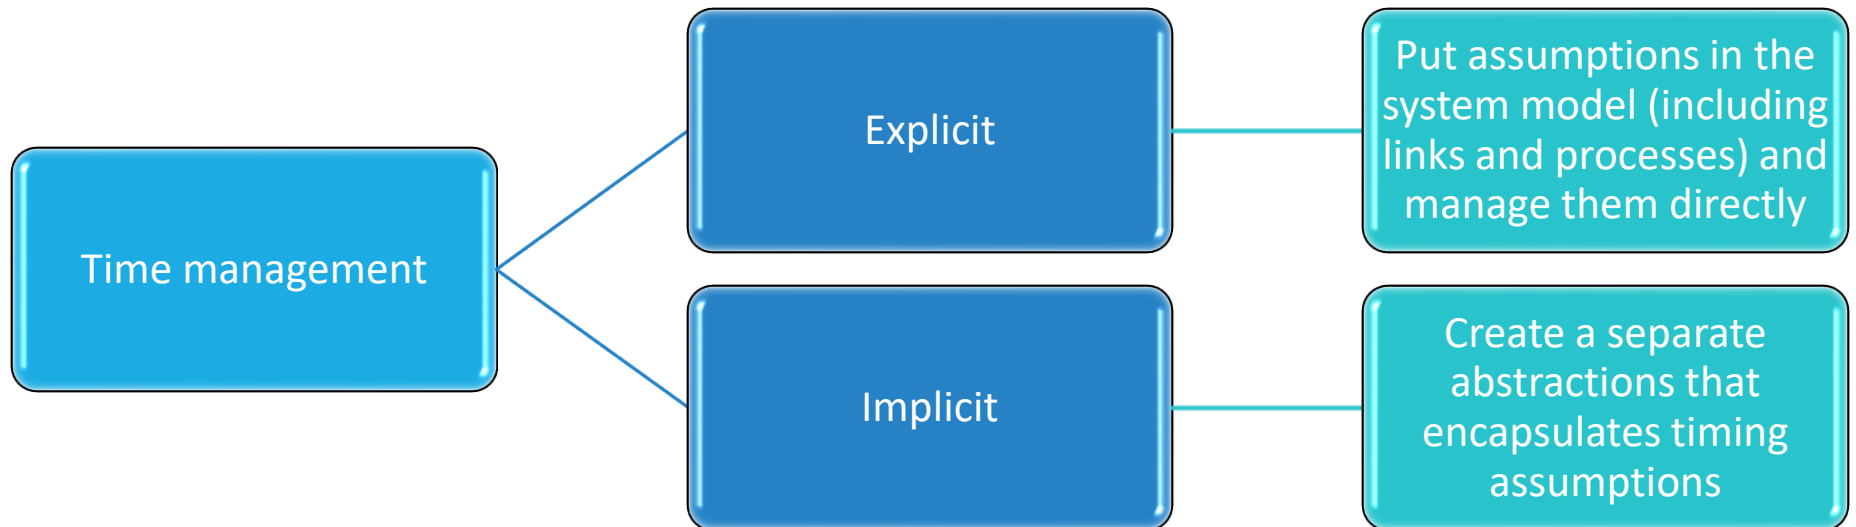
Partial synchrony

- requires abstract timing assumptions (after an unknown time t the system becomes synchronous)

Recap on Timing Assumptions

NOTE

- manipulating time inside a protocol/algorithm is complex and the correctness proof may become very involved and sometimes prone to errors



Failure Detector Abstraction

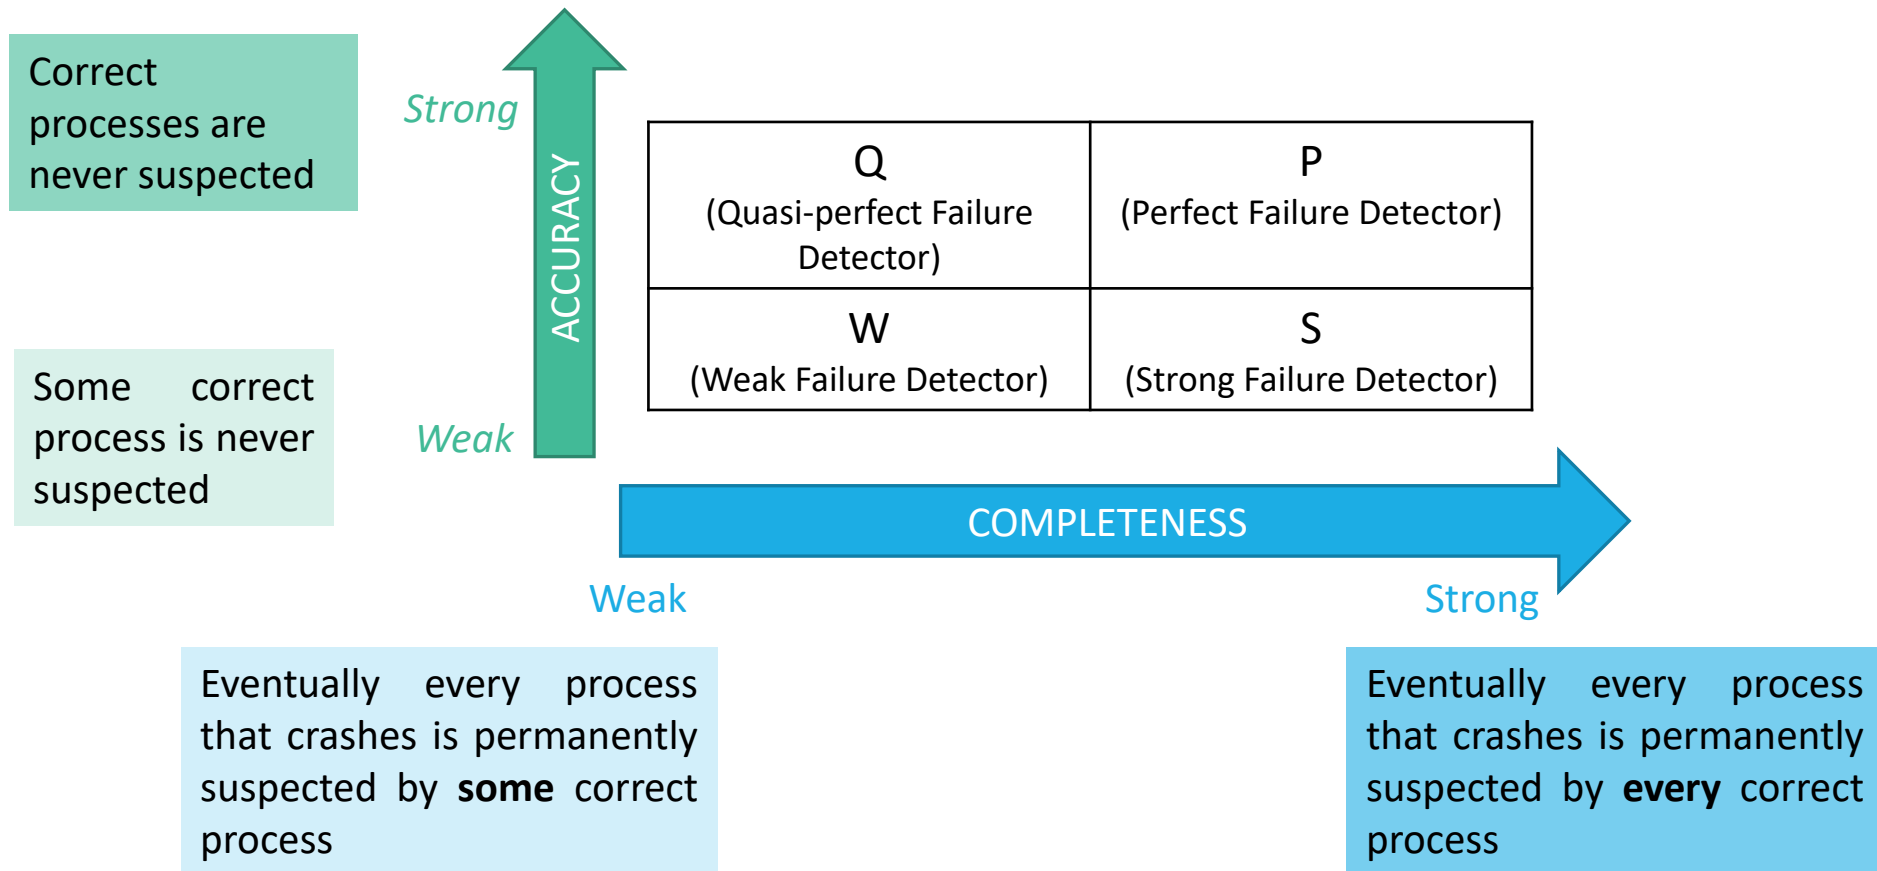
A Failure Detector is an oracle providing information about the failure state of a process

- it is a software module to be used together with process and link abstractions
- It encapsulates timing assumptions of either partially synchronous or fully synchronous system
 - Stronger are the timing assumption, more accurate the information provided by a failure detector will be

A failure detector is generally described by two properties:

- Accuracy (informally is the ability to avoid mistakes in the detection)
- Completeness (informally is ability to detect all failures)

Failure Detectors Classification



Failure Detectors Classification

OBSERVATION

- W guarantees that there is at least one correct process that is never suspected.
- Practically speaking, this could be difficult to achieve and thus it is worth to consider accuracy properties to hold eventually

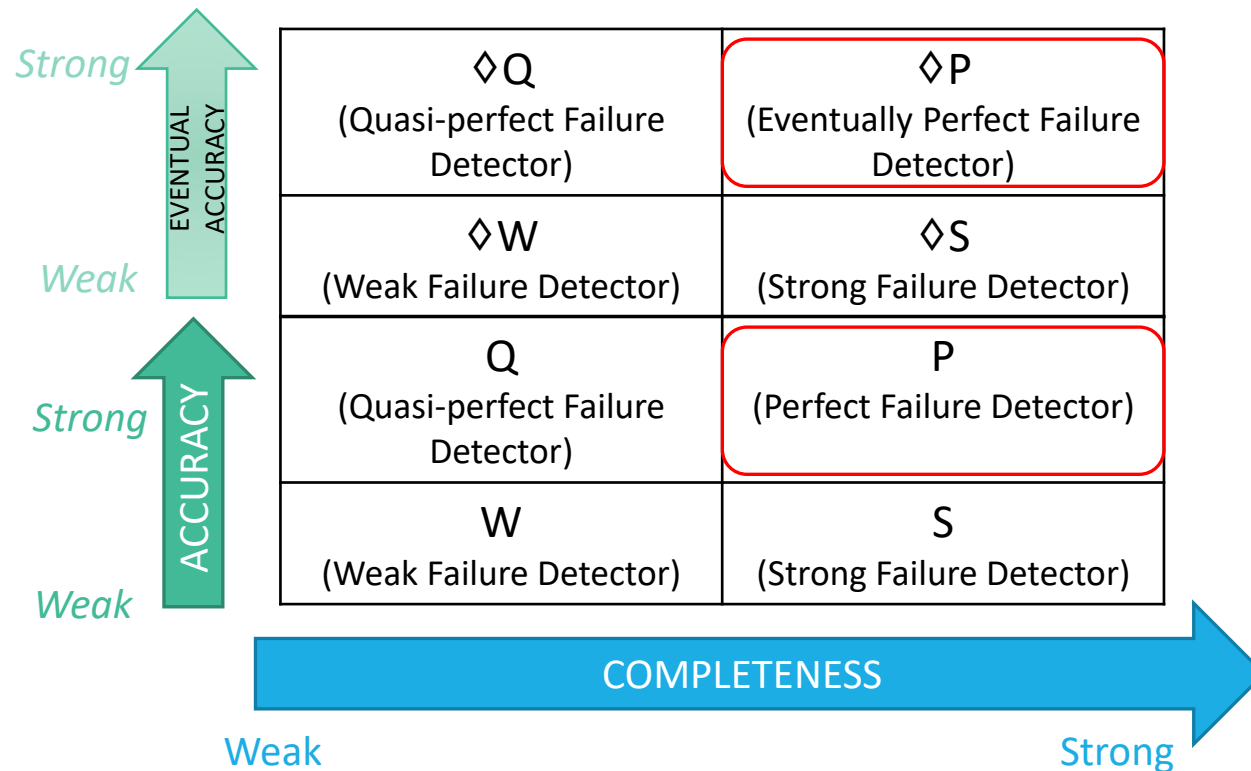
EVENTUAL STRONG ACCURACY

- There is a time after which correct processes are not suspected

EVENTUAL WEAK ACCURACY

- There is a time after which some correct process is not suspected

Failure Detectors Classification



Perfect Failure detectors (P)

System model

- synchronous system
- crash failures

Using its own clock and the bounds of the synchrony model, a process can infer if another process has crashed

Perfect failure detectors (P)

Specification

Module 2.6: Interface and properties of the perfect failure detector

Module:

Name: PerfectFailureDetector, **instance** \mathcal{P} .


Events:

Indication: $\langle \mathcal{P}, \text{Crash} \mid p \rangle$: Detects that process p has crashed.

Properties:

PFD1: *Strong completeness:* Eventually, every process that crashes is permanently detected by every correct process.

PFD2: *Strong accuracy:* If a process p is detected by any process, then p has crashed.



Perfect failure detectors (P) Implementation

Algorithm 2.5: Exclude on Timeout

Implements:

PerfectFailureDetector, **instance** \mathcal{P} .

Uses:

PerfectPointToPointLinks, **instance** pl .

upon event $\langle \mathcal{P}, \text{Init} \rangle$ **do**

$alive := \Pi$;
 $detected := \emptyset$;
 $starttimer(\Delta)$;

upon event $\langle \text{Timeout} \rangle$ **do**

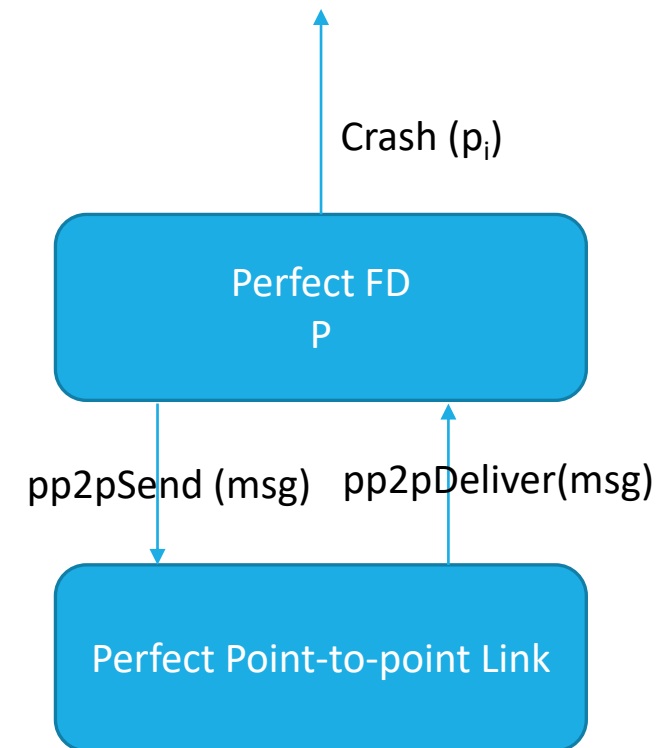
forall $p \in \Pi$ **do**
 if $(p \notin alive) \wedge (p \notin detected)$ **then**
 $detected := detected \cup \{p\}$;
 trigger $\langle \mathcal{P}, \text{Crash} \mid p \rangle$;
 trigger $\langle pl, \text{Send} \mid p, [\text{HEARTBEATREQUEST}] \rangle$;
 $alive := \emptyset$;
 $starttimer(\Delta)$;

upon event $\langle pl, \text{Deliver} \mid q, [\text{HEARTBEATREQUEST}] \rangle$ **do**

trigger $\langle pl, \text{Send} \mid q, [\text{HEARTBEATREPLY}] \rangle$;

upon event $\langle pl, \text{Deliver} \mid p, [\text{HEARTBEATREPLY}] \rangle$ **do**

$alive := alive \cup \{p\}$;



Correctness

- To prove the correctness, we must prove that both Strong Completeness and Strong Accuracy are satisfied
- What if links are fair loss?
- What if we select a timeout too long?
- What if we select a timeout too short?

Eventually perfect failure detectors ($\diamond P$)

System model

- partial synchrony
- Crash failures
- Perfect point-to-point links

Crashes can be accurately detected only after a (unknown) time t

- Before time t the systems behaves as an asynchronous one
- The failure detector may make mistake before time t considering correct processes as crashed.
- The notion of detection becomes suspicious

Eventually perfect failure detectors ($\diamond P$) Specification

Module 2.8: Interface and properties of the eventually perfect failure detector

Module:

Name: EventuallyPerfectFailureDetector, **instance** $\diamond P$.

Events:


Indication: $\langle \diamond P, \text{Suspect} \mid p \rangle$: Notifies that process p is suspected to have crashed.

Indication: $\langle \diamond P, \text{Restore} \mid p \rangle$: Notifies that process p is not suspected anymore.

Properties:

EPFD1: *Strong completeness:* Eventually, every process that crashes is permanently suspected by every correct process.

EPFD2: *Eventual strong accuracy:* Eventually, no correct process is suspected by any correct process.



Basic constructions rules of an eventually perfect FD

Use timeouts to suspect processes that did not sent expected messages

A suspect may be wrong

- A process p_i may suspect another one p_j as the current timeout is too short

◇ P is ready to change its judgment as soon as it receives a message from p_j

- In this case, the timeout value is updated

If p_j has actually crashed, p_i does not change its judgment anymore

Eventually perfect failure detectors ($\diamond P$) Implementation

Algorithm 2.7: Increasing Timeout

Implements:

EventuallyPerfectFailureDetector, **instance** $\diamond P$.

Uses:

PerfectPointToPointLinks, **instance** pl .

upon event $\langle \diamond P, \text{Init} \rangle$ do

$alive := \Pi$;
 $suspected := \emptyset$;
 $delay := \Delta$;
 $starttimer(delay)$;

upon event $\langle \text{Timeout} \rangle$ do

if $alive \cap suspected \neq \emptyset$ **then**
 $delay := delay + \Delta$;

forall $p \in \Pi$ do

if $(p \notin alive) \wedge (p \notin suspected)$ **then**
 $suspected := suspected \cup \{p\}$;
 trigger $\langle \diamond P, \text{Suspect} \mid p \rangle$;

else if $(p \in alive) \wedge (p \in suspected)$ **then**
 $suspected := suspected \setminus \{p\}$;
 trigger $\langle \diamond P, \text{Restore} \mid p \rangle$;

trigger $\langle pl, \text{Send} \mid p, [\text{HEARTBEATREQUEST}] \rangle$;

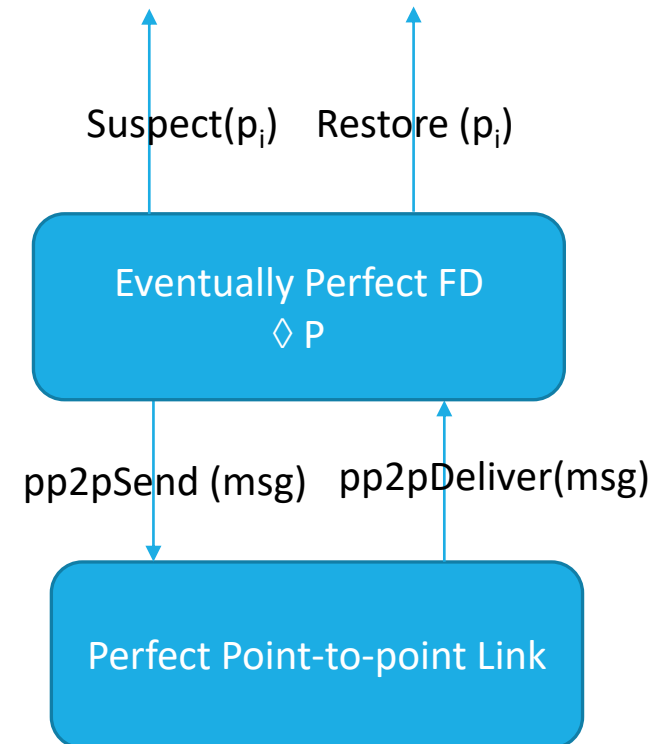
$alive := \emptyset$;
 $starttimer(delay)$;

upon event $\langle pl, \text{Deliver} \mid q, [\text{HEARTBEATREQUEST}] \rangle$ do

trigger $\langle pl, \text{Send} \mid q, [\text{HEARTBEATREPLY}] \rangle$;

upon event $\langle pl, \text{Deliver} \mid p, [\text{HEARTBEATREPLY}] \rangle$ do

$alive := alive \cup \{p\}$;



Correctness

Strong completeness

- If a process crashes, it will stop to send messages. Therefore the process will be suspected by any correct process and no process will revise the judgement.

Eventual strong accuracy

- After time T the system becomes synchronous. i.e., after that time a message sent by a correct process p to another one q will be delivered within a bounded time. If p was wrongly suspected by q , then q will revise its suspicious.

References

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 2 – from Section 2.6.1 to Section 2.6.5

T. Chandra, S. Toueg *Unreliable Failure Detectors for Reliable Distributed Systems*

- <https://dl.acm.org/doi/pdf/10.1145/226643.226647>