

Dependable Distributed Systems

Master of Science in Engineering in Computer Science

AA 2022/2023

LECTURES 23: CONSISTENCY CRITERIA FOR DISTRIBUTED
SHARED MEMORIES

Motivation

Distributed Shared Memories (DSMs) are an alternative to message passing for allowing inter-process communication

DSM Advantages

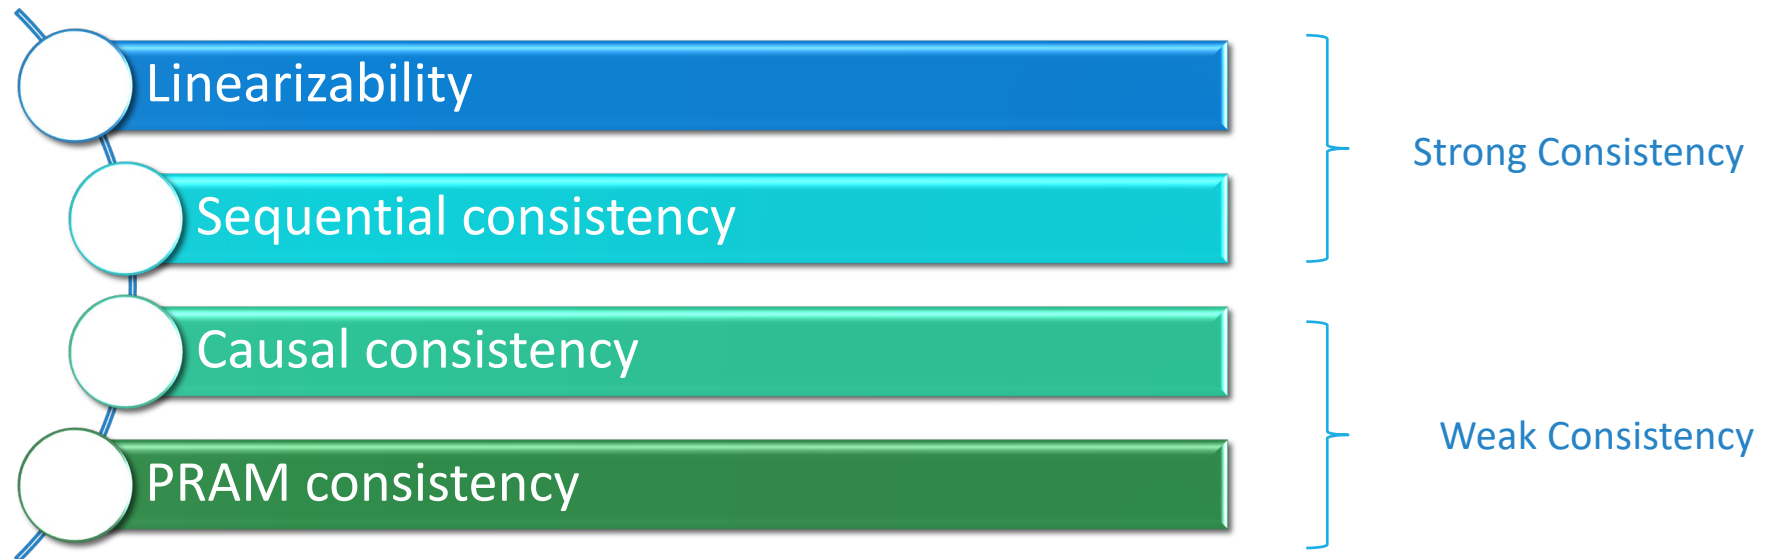
- support developers with the shared variables programming paradigm
 - abstract the underlying system
- increasing transparency with respect to portability, load balancing and process migration

The semantic of a shared memory is expressed by a consistency criterion

Consistency criteria

A consistency criterion defines the result returned by an operation

- It can be seen as a contract between the programmer and the system implementing replication



Notation

A shared memory system is composed by a set of sequential processes p_1, p_2, \dots, p_n interacting with a finite set of shared objects

Each object can be accessed by invoking read and write operations


- $w_i(x)v$ denote a write operation issued by p_i on the object x and writing the value v
- $r_j(x)v$ denote a read operation issued by p_j on the object x and returning the value v
- $op_i(x)v$ denote a generic operation issued by p_i on the object x and writing/returning the value v

Histories

Informally, an history represents the partial order of all the operations executed on the shared objects

The *Local history* \hat{h}_i is the sequence of operations issued by p_i

- if op1 and op2 are issued by p_i and op1 is issued first, than we say that op1 precedes op2 in p_i 's process order and we will denote it as $op1 \rightarrow_i op2$



OBSERVATION


- Given the set of operations h_i issued by p_i , the local history \hat{h}_i is the total order (h_i, \rightarrow_i)

Execution history

An execution history \hat{H} of a shared memory system is a partial order (H, \rightarrow_H) such that

- $H = \bigcup_i h_i$
- $op1 \rightarrow_H op2$ if:
 1. $op1$ and $op2$ are in *process-order* relation (i.e., there exists a p_i such that $op1 \rightarrow_i op2$) OR
 2. $op1$ and $op2$ are in *read-from-order* relation (i.e., $op1 = w_i(x)v$ and $op2 = r_j(x)v$) OR
 3. there exists $op3$ such that $op1 \rightarrow_H op3$ and $op3 \rightarrow_H op2$


If neither $op1 \rightarrow_H op2$ nor $op2 \rightarrow_H op1$ then $op1$ and $op2$ are said to be *concurrent*



Legality

A read operation $r(x)v$ is **legal** if

1. there exists a write $w(x)v$ preceding it in the history (i.e., $\exists w(x)v : w(x)x \rightarrow_H r(x)v$) AND
2. there not exists any other operation in between that write/read a different value u (i.e., $\nexists op(x)u : (u \neq v) \wedge w(x)v \rightarrow_H op(x)u \rightarrow_H r(x)v$)




A history is legal if all its read are legal

Sequential Consistency

A history $\hat{H} = (H, \rightarrow_H)$ is *sequentially consistent* if it admits a linear extension in which all the read are legal.

A *linear extension* $\hat{S} = (S, \rightarrow_S)$ of a partial order $\hat{H} = (H, \rightarrow_H)$ is a topological sort of this partial order (i.e., \hat{S} create a total order by maintaining the order of all ordered pairs in \hat{H})


Example



The execution is not linearizable!

$$\hat{S} = w1(x)0, w2(x)1, w2(y)2, r1(y)2, r2(x)1, r1(x)0$$

Example



**However, it is sequential
consistent**

$$\hat{S} = w2(x)1, w2(y)2, r2(x)1, w1(x)0, r1(y)2, r1(x)0$$

Linearizability vs Sequential Consistency

RECALL

- Linearizability requires that the linear extension \hat{S} also respects the real time of invocation
- Sequential consistency removes the real time aspect and just focus on logical time

Linearizability \Rightarrow Sequential Consistency

Sequential Consistency \nRightarrow Linearizability

Causal Consistency


Let $\hat{H} = (H, \rightarrow_H)$ be an execution history and let \hat{H}_i be the sub-history of \hat{H} from which all the read operations not issued by p_i are removed.

A history $\hat{H} = (H, \rightarrow_H)$ is *causally consistent* if, for each process p_i , all the read operations of \hat{H}_i are legal.

OBSERVATION

- in a causally consistent history, all processes see the same partial order of operations but each process may see a different linear extension.


Example



The execution is not sequential consistent!

$\hat{S} = w1(x)0, w2(y)0, w2(y)2, r2(x)0, w1(x)1, r2(x)1, r1(y)0, r1(y)2$

Example




**However, it is causal
consistent**

$$\hat{S}_1 = w1(x)0, w1(x)1, w2(y)0, r1(y)0, w2(y)2, r1(y)2$$

$$\hat{S}_2 = w2(y)0, r1(y)0, w1(x)0, r2(x)0, w1(x)1, r2(x)1$$

Example



Sequential Consistency \Rightarrow Causal Consistency
Causal Consistency \nRightarrow Sequential Consistency

References

Michel Raynal and André Schiper: “A suite of formal definitions for consistency criteria in distributed shared memories”

available at:

<https://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.52.6880&rep=rep1&type=pdf>