

Dependable Distributed Systems

Master of Science in Engineering in Computer Science

AA 2022/2023

LECTURE 30,31: DLT AND BLOCKCHAIN

What is a Distributed Ledger?

Distributed Ledger

A ledger is a written or computerized record of all the transactions a business has completed




A distributed ledger is a database replicated across several locations or among multiple participants used to store record of transactions


What is a Distributed Ledger?

Key properties

- Consistency
- Integrity
- Availability

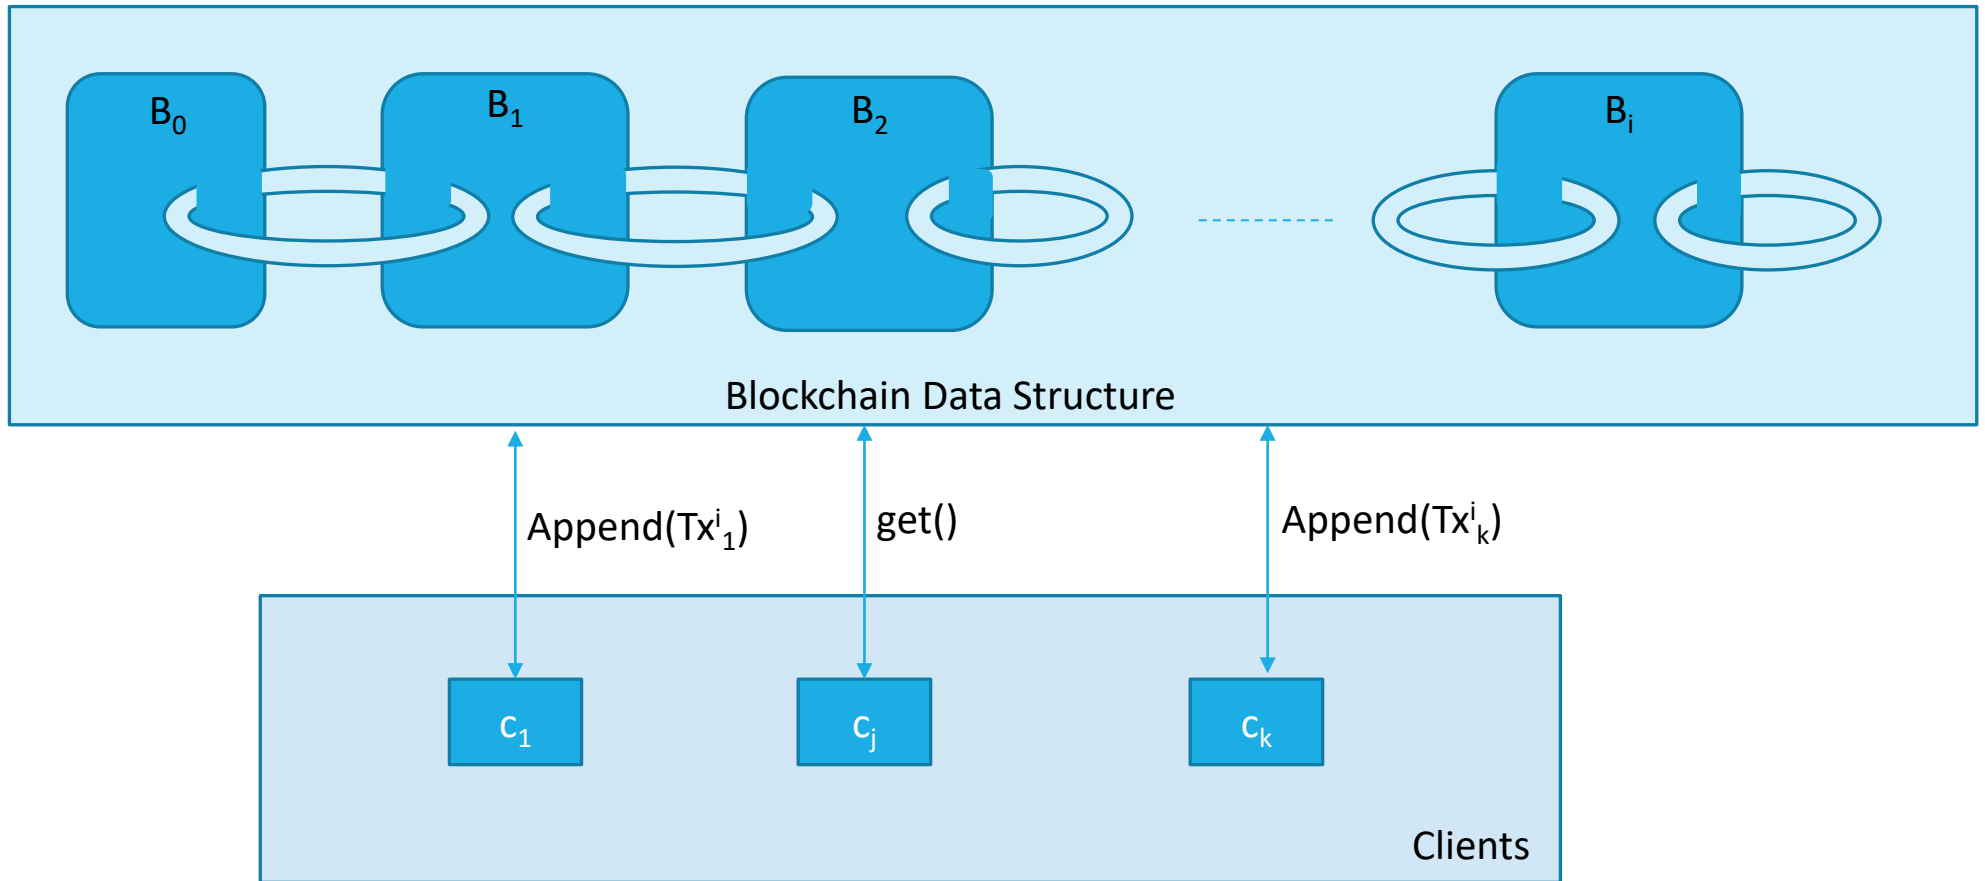


What is a blockchain?



A blockchain is a **decentralized, distributed** data structure used to **record transactions** (aggregated in blocks) across many computers.


What is a blockchain?



What is a Blockchain?

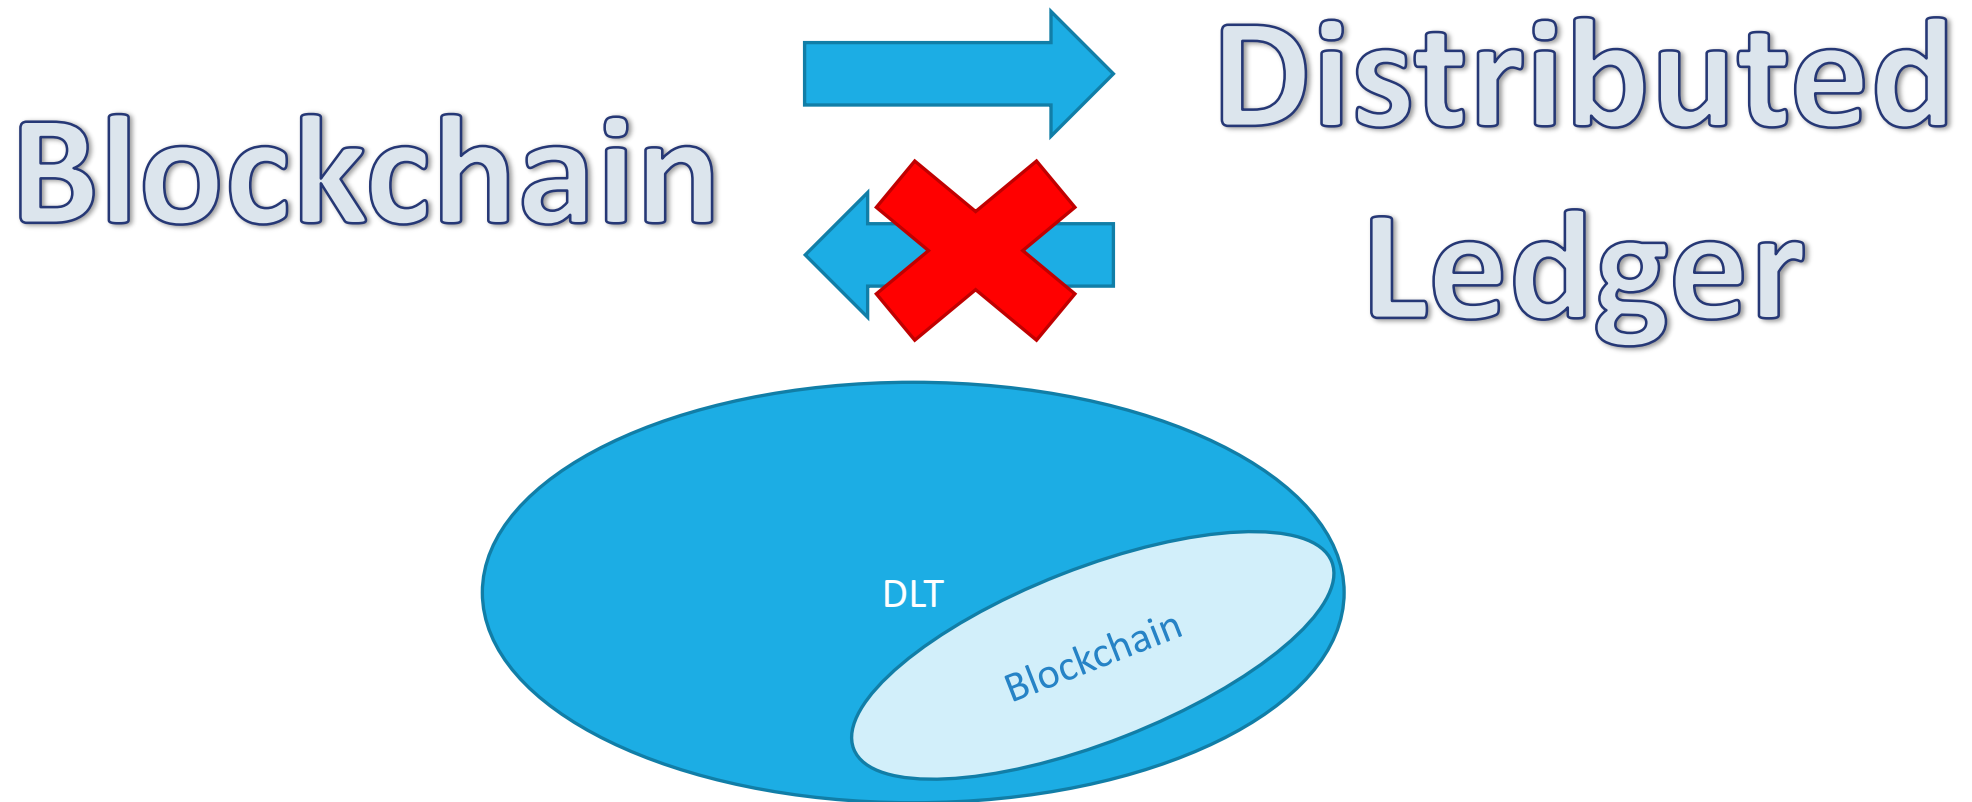
Key properties

- Consistency
- Integrity
- Availability



Blockchain vs DL


Every blockchain represents a distributed ledger but the vice versa is not true



What is a transaction?

A transaction is


- an instance of buying or selling something
- an exchange or interaction between people



What is a Block?

Every block in a blockchain contains a set of transaction

- Clients generate transactions and submit them to nodes implementing the blockchain



How to create a Block?


Transactions are collected by processes

Transactions need to be validated i.e., they need to be valid with respect to the ledger specification

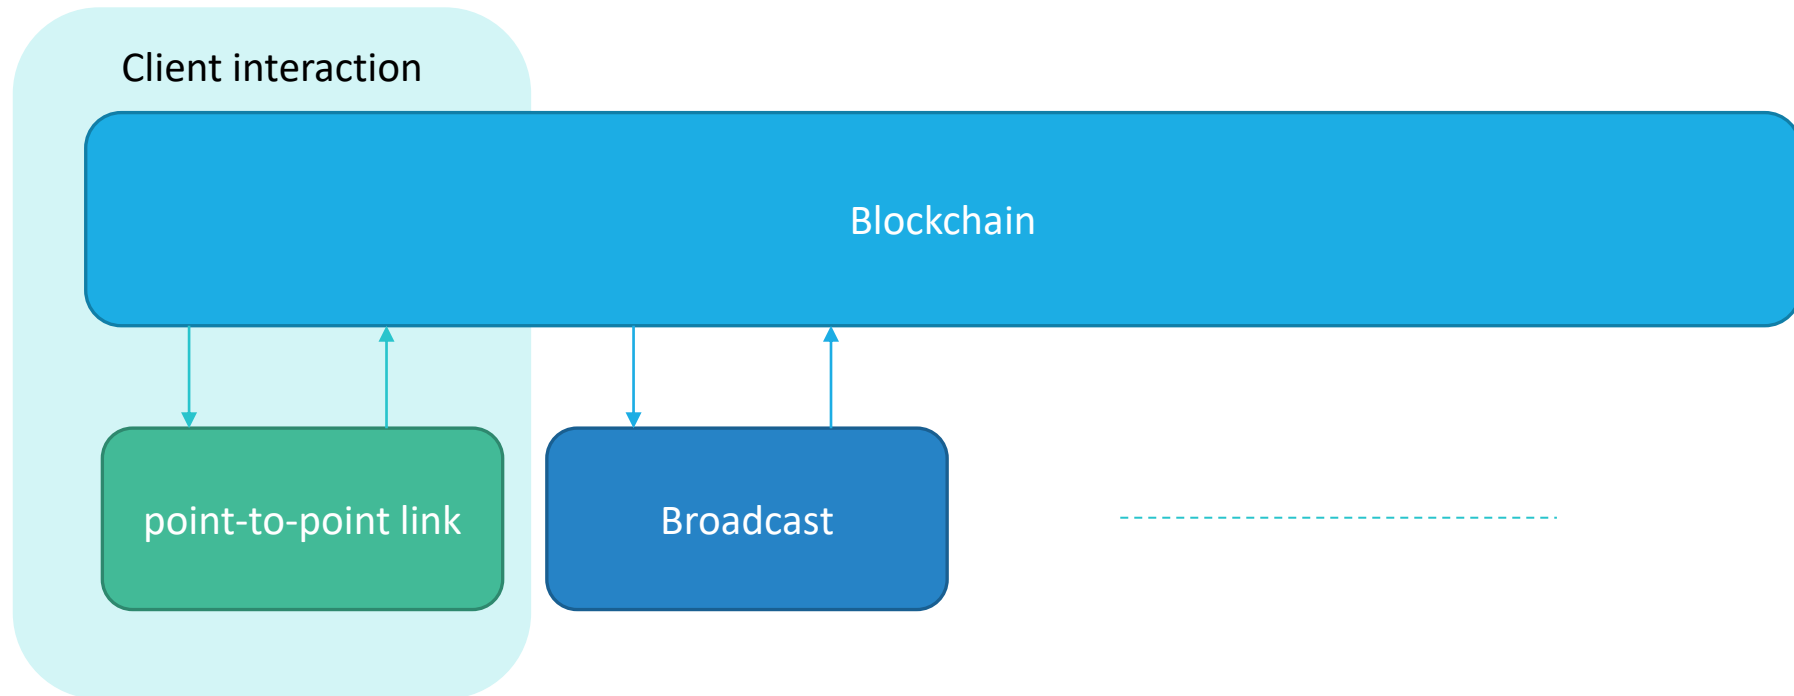
- E.g., in a financial ledger a transaction cannot spend money that are not in the account

Transactions remain *unconfirmed* until they are inserted in a block that is successfully attached to the chain

When "enough" transactions have been collected a block can be created and attached



Key ingredients



How to attach a block to the chain?

Validate transaction

- In order to validate a transaction, nodes need to read the “last” state of the ledger and check that the current transaction is correct with respect to the semantics of the ledger
 - E.g. if in the last block we found that Alice’s account has 300\$, then any transaction spending at most 300\$ will be valid


OBSERVATION

- Ordering of transactions in the blocks is fundamental to check and guarantee the validity of the ledger

How to attach a block to the chain?

Chaining blocks

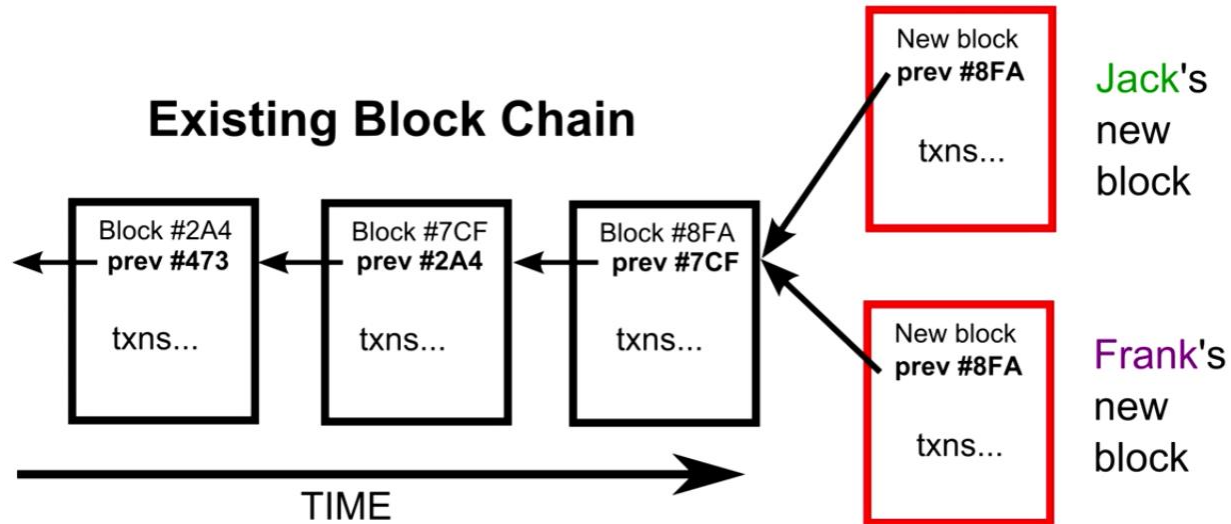
1. Check that all the transactions in the block are valid
2. Compute the hash $H()$ of the last block attached to the chain
3. Include this hash in the current block to generate a pointer and attach the current block



How to attach a block to the chain?

ISSUES

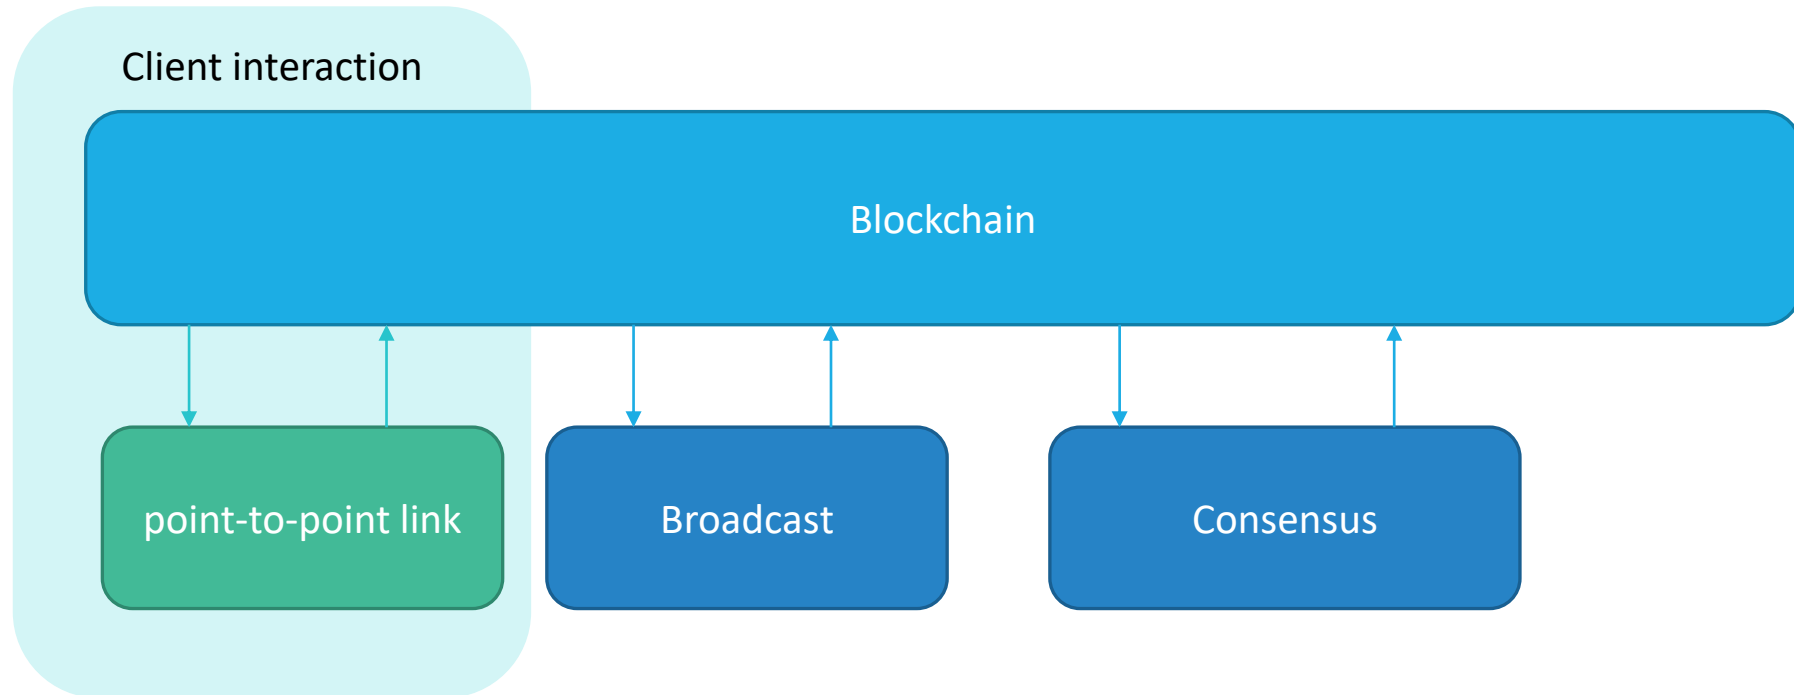
- Concurrency in the block creation process
- Asynchronous systems



How to create and chain blocks

**Attaching a new block
requires agreement
between nodes**

Key ingredients



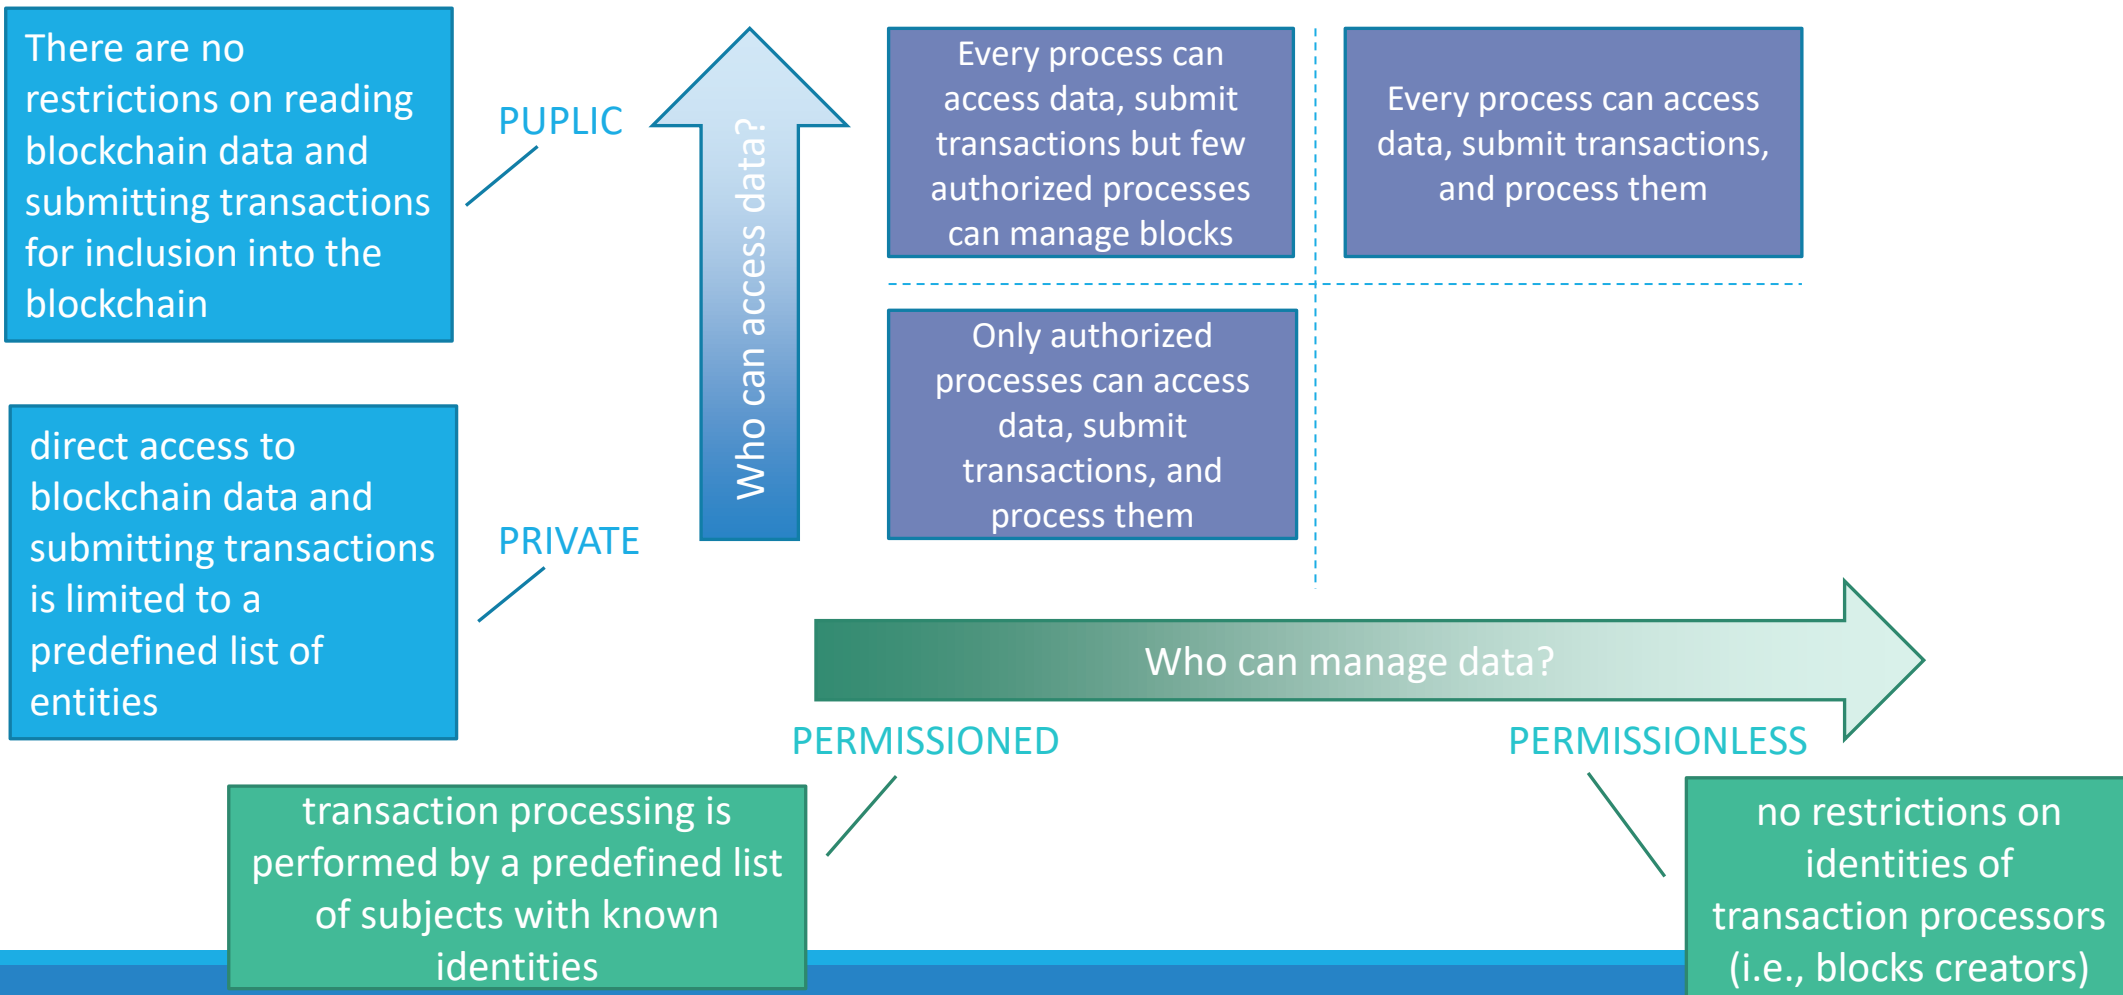
How to create and chain blocks

**Attaching a new block
requires agreement
between nodes**



**Which Consensus
protocol should I use?**

Blockchain Classification



How to create and chain blocks

Attaching a new block requires Consensus



**Which Consensus
protocol should I use?**

Public
Permissionless
Blockchain



Proof of X
(i.e., leader based
consensus)

OR

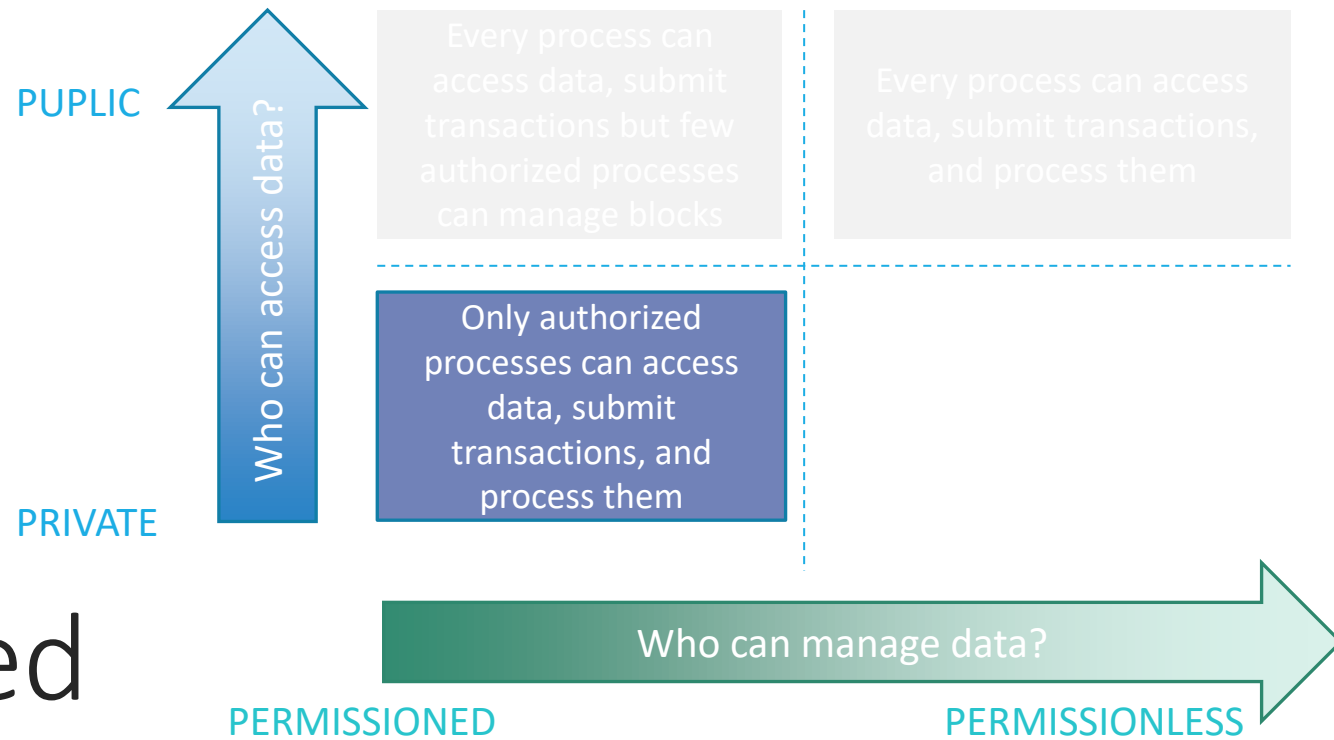
Private
Permissioned
Blockchain



BFT Consensus
(fully distributed consensus)

Private Permissioned Blockchain

THE PBFT PROTOCOL



Practical Byzantine Fault Tolerance

In private permissioned blockchain we have that

- Transactions are submitted only by known clients and only known clients can access them
- Blocks are created and attached by known and authorized processes


IDEA

- Processes maintain a copy of the Blockchain locally and run a BFT consensus protocol to agree on the next block
- PBFT is currently the adopted protocol
 - It has been designed to support replication
 - It merges the primary-backup approach with the Byzantine General Consensus approach

Practical Byzantine Fault Tolerance

System model:

- Asynchronous distributed system
 - nodes connected by network
 - messages can be lost, delayed, duplicated, no order
- Independent node failures (byzantine)



Basically the same assumed to solve the Byzantine Generals Problem

Practical Byzantine Fault Tolerance

System model:

- Cryptographic techniques (public key signatures, MAC, message digest produced by collision resistant hash functions)
 - needed to avoid spoofing, replay attacks and corruption
- All replicas know each other's public keys to verify signatures

Practical Byzantine Fault Tolerance

System model:

- Strong adversary that can
 - coordinate the action of faulty nodes
 - delay communication
 - delay correct nodes (not indefinitely)
- The adversary cannot subvert the cryptographic techniques cited above

Practical Byzantine Fault Tolerance

Properties:

- **Safety**
 - satisfies linearizability – like a centralized system that executes operations atomically one at a time
 - regardless of faulty clients
 - but faulty clients can do strange (legal) operations (e.g., write garbage in a file system)
 - limit damage by access control mechanisms

Practical Byzantine Fault Tolerance

Properties:

- Liveness
 - clients receive responses to their requests if at most $(n-1)/3$ replicas are faulty and $\text{delay}(t)$ does not grow faster than t indefinitely

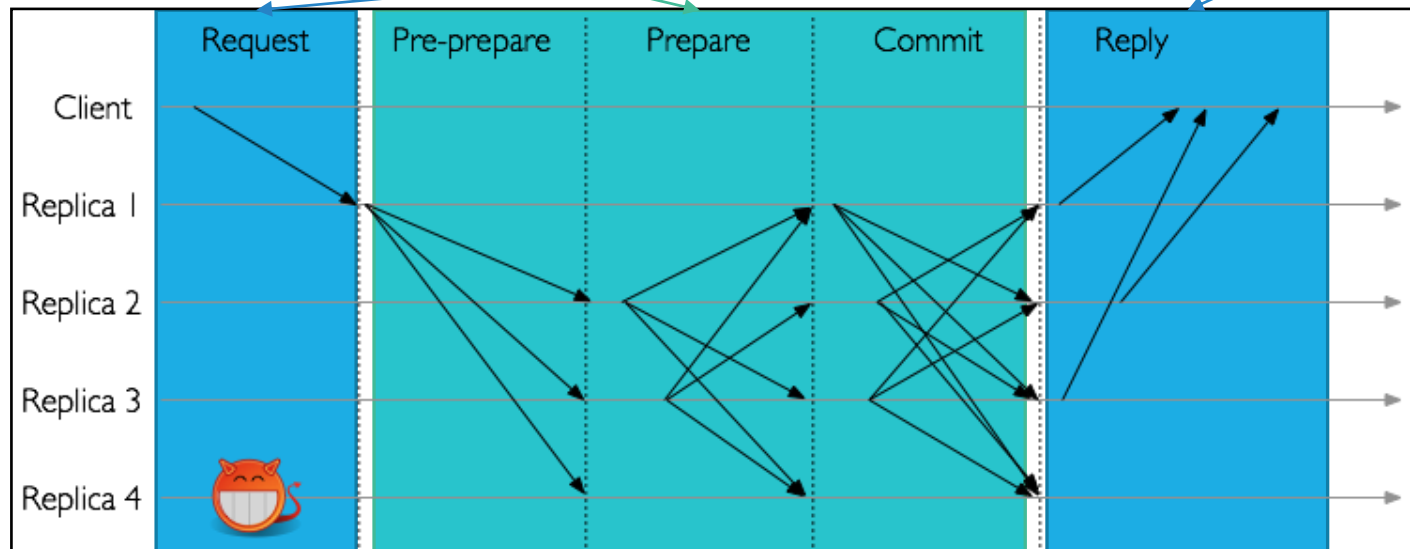
Practical Byzantine Fault Tolerance

The protocol is divided in five phases:

- Request, pre-prepare, prepare, commit and reply

Derived from Byzantine Generals

Derived from Primary-backup approach



Practical Byzantine Fault Tolerance

Basic idea:

- A client sends a request to invoke an operation to the primary
- The primary multicasts the request to the backups
- Replicas execute the request and send a reply to the client
- The client waits for $f+1$ replies from different replicas with the same result; this is the result of the operation

Practical Byzantine Fault Tolerance

If the primary is not correct it is substituted with a new one as soon as a misbehavior is detected

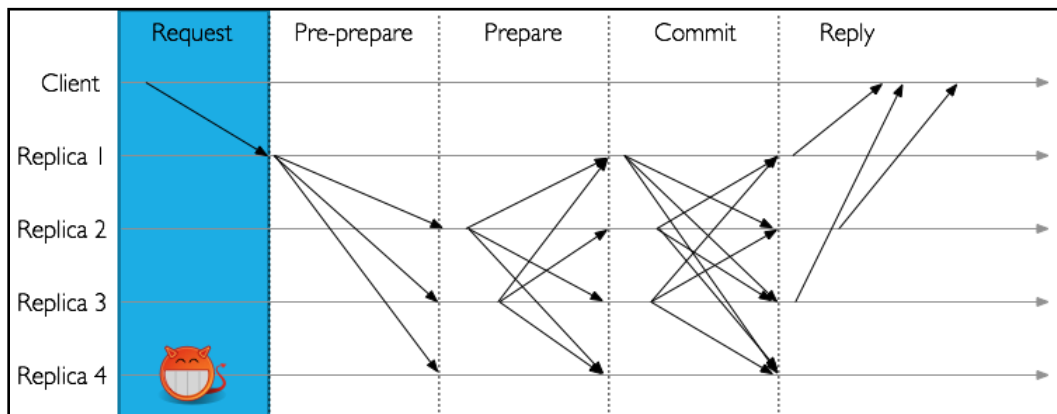
The system lifetime is characterized by a series of *views* each with a different primary

Views are numbered

Practical Byzantine Fault Tolerance

Request phase:

- Client c requests the execution of operation o issuing the request $\langle \text{REQUEST}, o, t, c \rangle_{sc}$ to the primary
- The request is sent to the last primary known by the client
 - If a response is not received by a timeout the request is multicast to all replicas
- t is a timestamp used to guarantee *exactly-once* semantic



Practical Byzantine Fault Tolerance

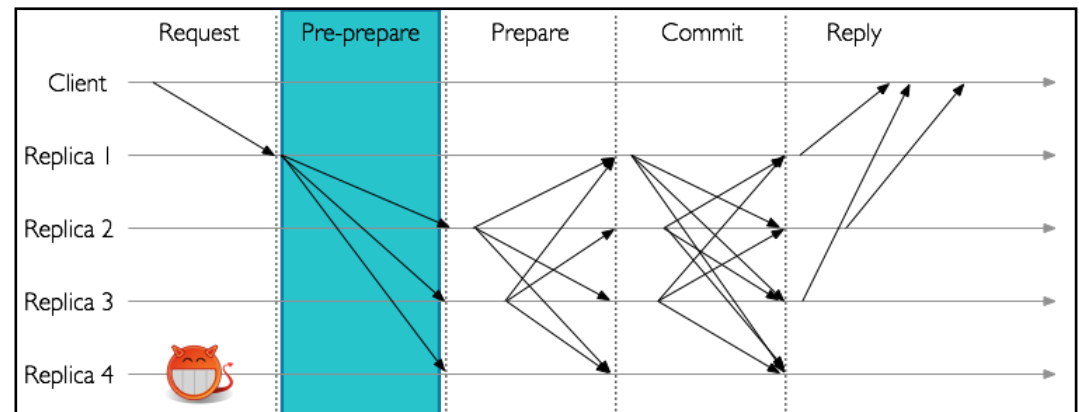
Replica internal state:

- replica id i (between 0 and $N-1$)
- service state
- view number v (initially 0)
- log of accepted messages

Practical Byzantine Fault Tolerance

Pre-prepare phase:


- the primary p sends to all backups the message $\langle \text{PRE_PREPARE}, c, n, d \rangle_{\text{op}}, m \rangle$
- c is the client id
- n is a sequence number for the current view assigned by the primary
- d is a digest of the client request
- m is the client request



Practical Byzantine Fault Tolerance

Pre-prepare phase:

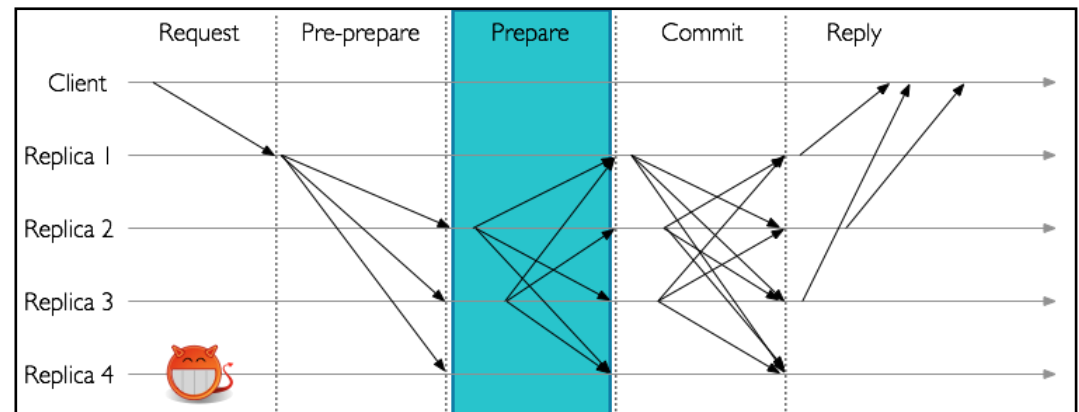
- a backup accepts (puts in the log) the pre-prepare msg iff:
 1. the signature in the pre-prepare and request messages are verified, and d is m 's digest
 2. it is currently in view v
 3. it has not accepted a pre-prepare message for view v and sequence number n containing a different digest
 4. the sequence number n is between a low watermark h and a high watermark H



Practical Byzantine Fault Tolerance

Prepare phase:

- a backup i sends to all nodes the message $\langle \text{PREPARE}, v, n, d, i \rangle_{\sigma_i}$
- another node accepts a prepare message (and puts it in the log) iff:
 1. signatures are correct
 2. it is currently in view v
 3. n is between h and H



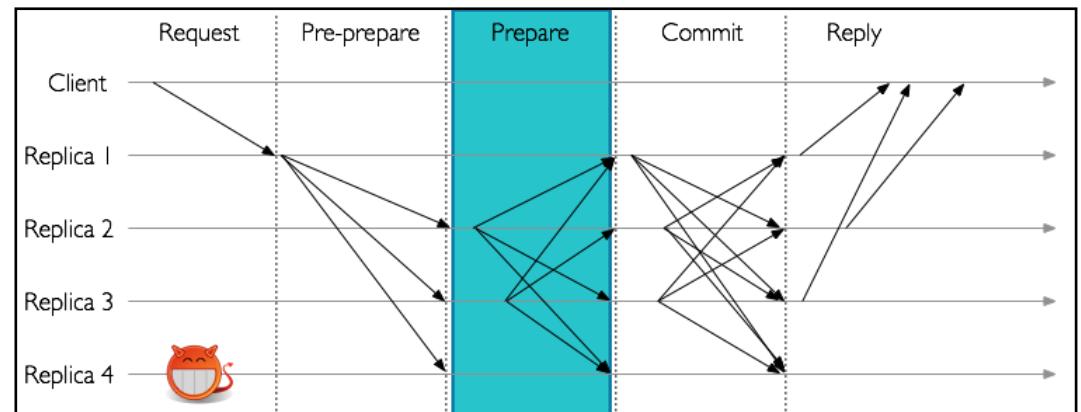
Practical Byzantine Fault Tolerance

Replicas remain blocked until the predicate $prepared(m, v, n, i)$ becomes true

This happen iff replica i has in its log:

1. the message m
2. a pre-prepare for m
3. More than $2f$ prepare messages from different backups that match the pre-prepare

matching among messages is done checking v , n and d



Practical Byzantine Fault Tolerance

The pre-prepare and prepare phase guarantee that non-faulty replicas agree on a total order of execution for requests in a same view


More precisely, they ensure that if $prepared(m, v, n, i)$ is true then $prepared(m', v, n, j)$ is false for any non faulty replica j and any m' such that the digest of m is different from the digest of m'

This is due to the intersection among quorums of nodes that accept the prepare messages

Practical Byzantine Fault Tolerance

Commit phase:

- as soon as $prepared(m, v, n, i)$ becomes true replica i sends to all the message $\langle \text{COMMIT}, v, n, d, i \rangle_{\sigma_i}$
- another node accepts a commit message (and puts it in the log) iff:
 1. signatures are correct
 2. it is in view v
 3. n is between h and H



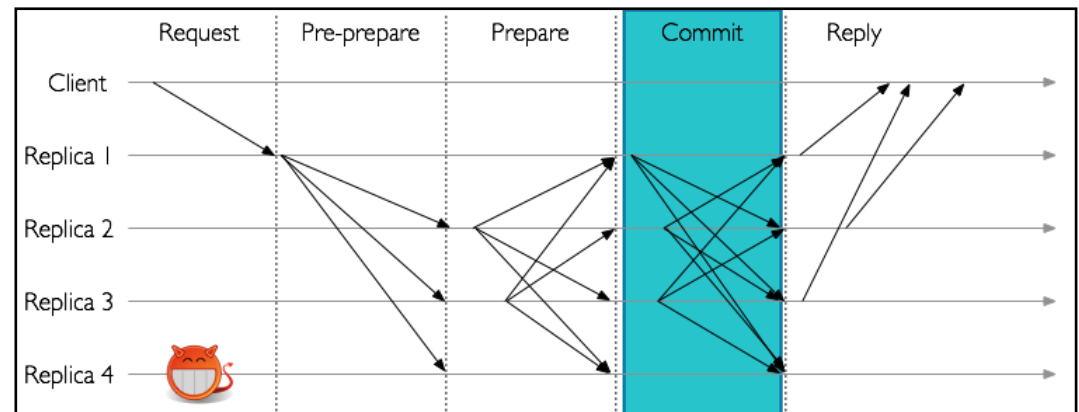
Practical Byzantine Fault Tolerance

The predicate $committed(m, v, n)$ is true iff $prepared(m, v, n, i)$ is true for all i in a set of $f+1$ non faulty replicas.

The predicate $committed-local(m, v, n, i)$ is true iff $prepared(m, v, n, i)$ is true and i has accepted $2f+1$ commits from different replicas that match the pre-prepare for m

The commit phase guarantees that if $committed-local(m, v, n, i)$ is true for some non faulty i , then $committed(m, v, n)$ is true.

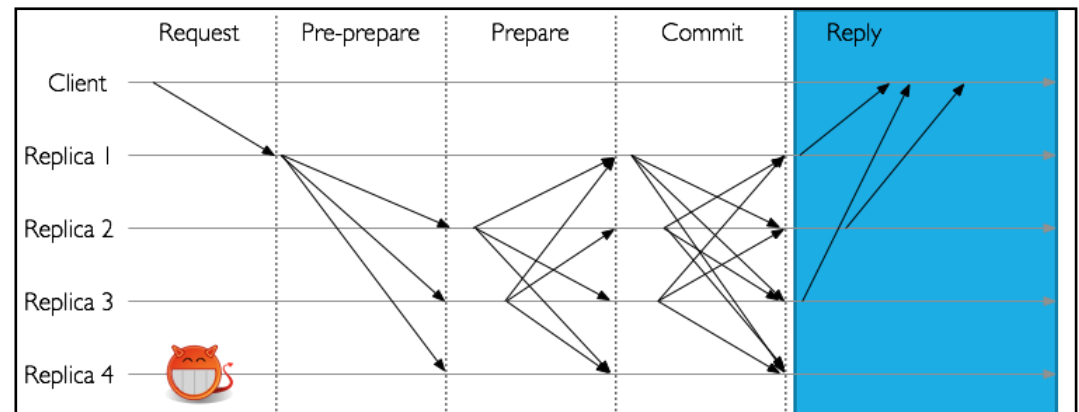
- because if i received $2f+1$ commits, $prepared(m, v, n, i)$ is true for at least $f+1$ non-faulty nodes.



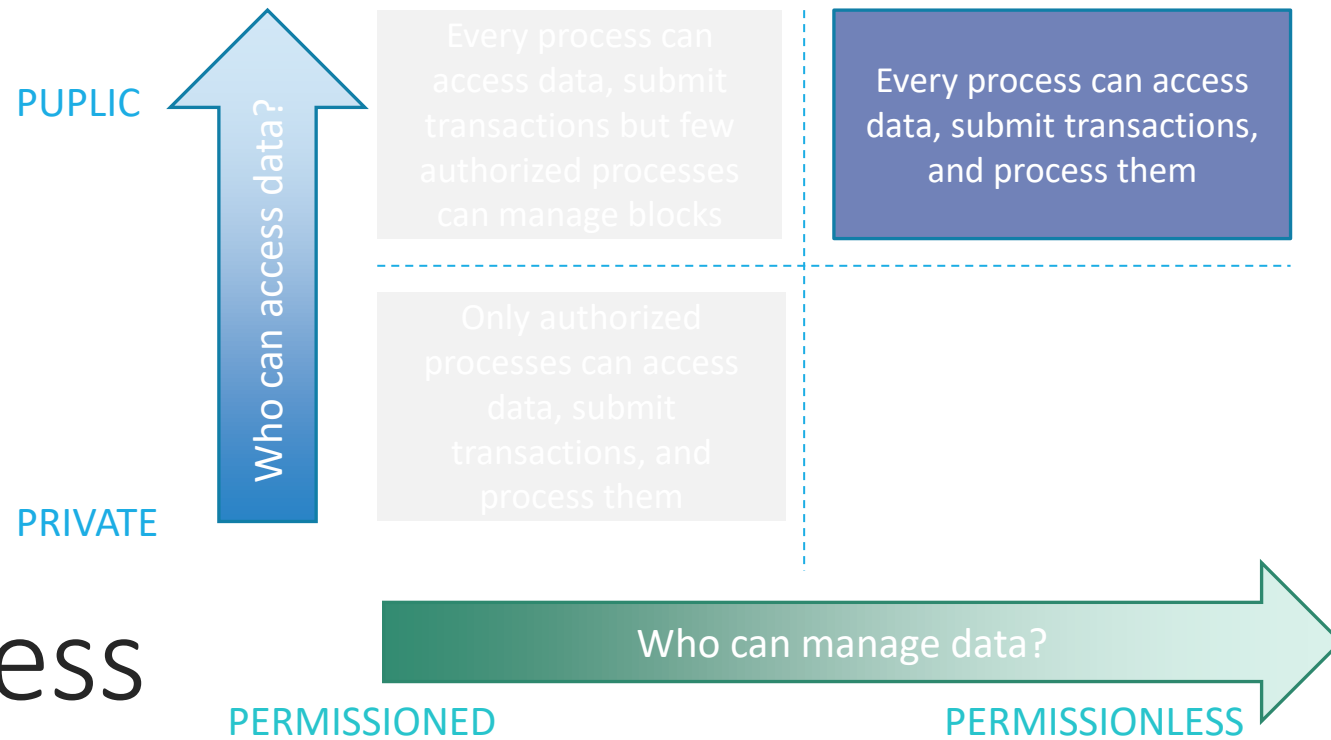
Practical Byzantine Fault Tolerance

Reply phase:

- A replica i executes the request as soon as
 - $committed-local(m, v, n, i)$ is true
 - i 's state reflects the sequential execution of all requests
- The response is then sent to the client $\langle \text{REPLY}, v, t, c, i, r \rangle_{\sigma_i}$
- The client waits for $f+1$ replies with valid signatures and the same values for t and r



Public Permissionless Blockchain



PROOF OF X

Public Permissionless Blockchain

In public permissionless blockchain we have that

- Transactions can be submitted by everyone and every process can read the chain
- Blocks can be created and attached by every process in the system

OBSERVATION

- Public permissioned blockchains are characterized by
 - Lack of trust in other processes
 - Possibly large scale

CONSEQUENCE

- PBFT-like protocols do not work

Public Permissionless Blockchain

IDEA

- Processes start a competition and only the winner can attach its block to the blockchain (i.e., they are trying to elect a leader)
- They implement a “randomized leader election”

ISSUES


- Two processes may win the competition



- Two blocks can be attached to the chain


Proof-of-work (PoW) & Mining

- **Proof-of-work:** mathematical challenge to “solve” a block
- **Mining:** find a number s.t. $\text{hash}(\text{block}) < \text{target}$
- The first node who “solves” a **block** can propose it as next in the blockchain
- Other nodes which receive the block re-compute the hash to check the validity




Proof-of-work & Mining: Branches Management

- Occasionally two or more blocks may arrive together → **temporal disagree**
 - The probability that two or more nodes mine a block at the same time is very low
- **RULE:** in case of branches the network has to converge to the longest branch.




Proof-of-work & Mining: Branches Management

- Occasionally two or more blocks may arrive together → **temporal disagree**
 - The probability that two or more nodes mine a block at the same time is very low
- **RULE:** in case of branches the network has to converge to the longest branch.
 - **Problem solved with next blocks:** *eventually one branch will become the longest (usually after one block) bringing convergence*




Scalability Issues

- Transaction rate depends on two parameters
 - **Block size:** how many transactions to include in a block?
 - **Block interval:** how long to wait for a block to propagate to all the nodes?
 - Change run-time the **difficult of the target** to solve a block according to the computational power of the network
- Metrics
 - **Throughput:** how many transactions per second?
 - **Latency:** how long to wait for a transaction to complete?
- How parameters impact on metrics
 - Increasing block size improves throughput, but the resulting bigger blocks take longer to propagate in the network
 - Reducing the block interval reduces latency, but leads to instability where the system is in disagreement and the blockchain is subject to reorganization



Miner Requirements

- A miner needs a **high computational power** to compute a thousands of hash per seconds (KH/s or MH/s)
- CPUs are ineffective (few cores) better using GPUs



Miner Requirements

- Possible to improve performances by employing a cluster of GPU



KADA 6.1 GPU Mining Rig Open Air Frame Case Chassis with 6 USB Risers - Ethereum

99.6 % Positive feedback (★ ★ ★ ★ ★)

235 \$

Listing Status: **Completed**

Country: **US**

Item condition: **New**

Buy Now!

ebay

PayPal

VISA

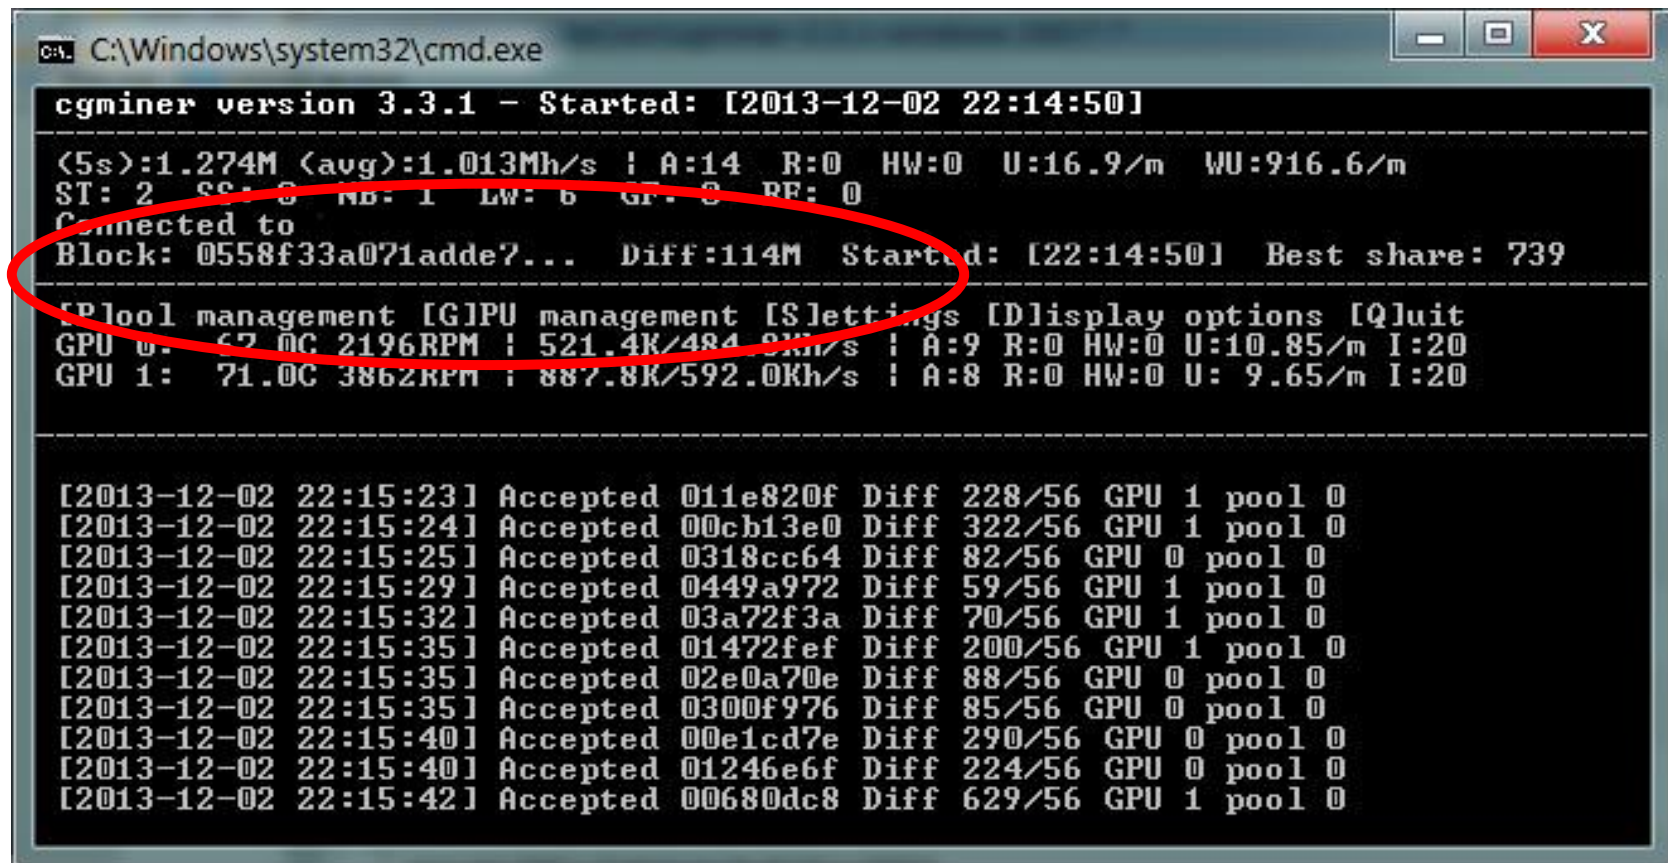
MasterCard

Discover

DISCOVER

Software for Mining

- CGminer, BFGminer, BitMiner, BTCMiner, DiabloMiner, ...



```
C:\Windows\system32\cmd.exe
cgminer version 3.3.1 - Started: [2013-12-02 22:14:50]
-----
<5s>:1.274M <avg>:1.013Mh/s | A:14 R:0 HW:0 U:16.9/m WU:916.6/m
ST: 2 SS: 0 NB: 1 LW: 6 GF: 0 RF: 0
Connected to
Block: 0558f33a071adde7... Diff:114M Started: [22:14:50] Best share: 739
-----
[P]ool management [G]PU management [S]ettings [D]isplay options [Q]uit
GPU 0: 62.0C 2196RPM | 521.4K/484.0Kh/s | A:9 R:0 HW:0 U:10.85/m I:20
GPU 1: 71.0C 3862RPM | 887.8K/592.0Kh/s | A:8 R:0 HW:0 U: 9.65/m I:20
-----
[2013-12-02 22:15:23] Accepted 011e820f Diff 228/56 GPU 1 pool 0
[2013-12-02 22:15:24] Accepted 00cb13e0 Diff 322/56 GPU 1 pool 0
[2013-12-02 22:15:25] Accepted 0318cc64 Diff 82/56 GPU 0 pool 0
[2013-12-02 22:15:29] Accepted 0449a972 Diff 59/56 GPU 1 pool 0
[2013-12-02 22:15:32] Accepted 03a72f3a Diff 70/56 GPU 1 pool 0
[2013-12-02 22:15:35] Accepted 01472fef Diff 200/56 GPU 1 pool 0
[2013-12-02 22:15:35] Accepted 02e0a70e Diff 88/56 GPU 0 pool 0
[2013-12-02 22:15:35] Accepted 0300f976 Diff 85/56 GPU 0 pool 0
[2013-12-02 22:15:40] Accepted 00e1cd7e Diff 290/56 GPU 0 pool 0
[2013-12-02 22:15:40] Accepted 01246e6f Diff 224/56 GPU 0 pool 0
[2013-12-02 22:15:42] Accepted 00680dc8 Diff 629/56 GPU 1 pool 0
```

ASIC

- ASIC: alternative specific hardware for mining
- Example: AntMiner S5
 - Hashrate: 1.15GH/s
 - Price: \$370
 - Power consumption: 590W
 - Mining for 24 hours/day is expensive!

Question:
Why should a user mine?



Miner Incentive

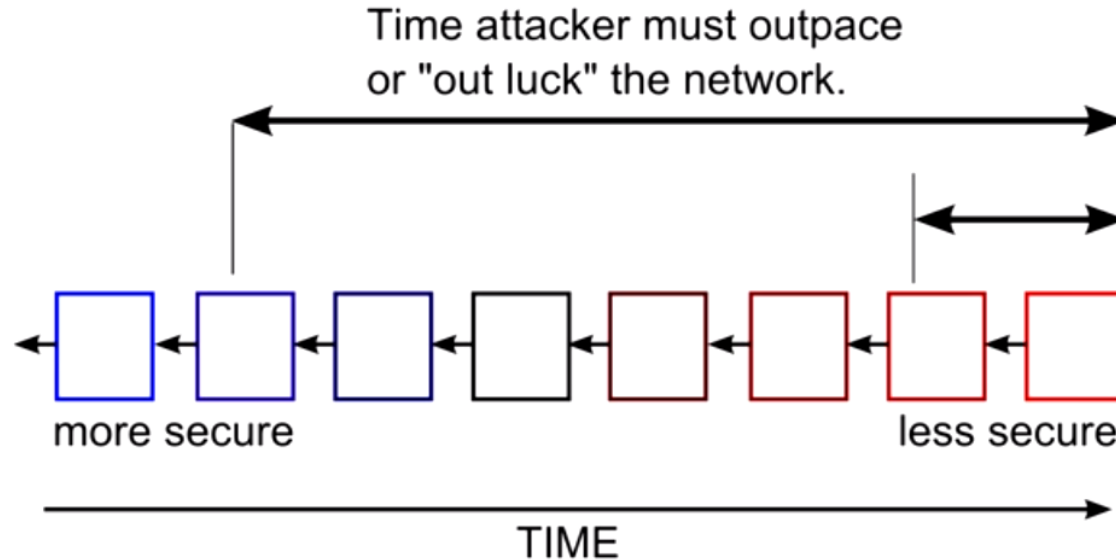
- **Reward:** Solving a block gives coins to node that found the proof-of-work
 - Some txn may bring **additional fee** to node who mines the block containing that txns
 - Currently the txn fees are quasi-zero
 - Miners will be motivated to include in a block txn with a fee
- Rewards are an incentive for nodes to keep them supporting the blockchain and keep nodes honest
 - Invalid txns won't give to miner the reward!
- Furthermore, it is a way to distribute coins into circulation
 - There is no central authority issuing new coins
 - Each crypto-currency platform will not erode new coins to miners forever → currency deflation
 - Is effective to mine whilst you earn more money than those spent for electric power

How the PoW Ensures Integrity: Computational Power

- **Impossible to change a txn in a block b in the blockchain:** an attacker should be quicker than the rest of the whole network to mine a block.
- In that case he could be able to re-mine $n+1$ blocks (i.e. all blocks next to $b + 1$) quicker than the rest of the network to mine a new block.
- If so, the attacker could obtain the longest (modified) blockchain and all network would converge to it.
- But for doing that, he would have the 50% of the computational power of the network to have a 50% probability to solve a block before another node.
- And he would have a higher percentage of computational power to solve more blocks sequentially.

How Blockchain Ensures Integrity: Computational Power

- Last blocks are so less secure
- Wait for 5/6 blocks makes a success probability too low for an attacker
- This solution protects from both **Integrity** and **Double Spending Fraud**



Alternative to PoW?

- **PoW pro:** very secure
- **PoW cons:** waste of electric power
- **Proof-of-Stake (PoS)** is the PoW alternative
 - Secure without mining → no energy wasted

Proof-of-Stake

- Instead of mine a block, the creator of the next block is chosen in a deterministic way according to its wealth (i.e. stake)
- The reward are not related to the created block but according to your wallet
- The longer you keep the coin in the wallet, the more the reward is high
- The probability to ***mint*** (instead of mine) a block is proportional to your wallet
 - **Minting** process require a lot of coin to attack the network
 - If you have a the $p\%$ of coins of the network you will mint the $p\%$ of the blocks
 - Very difficult to mint two consecutive blocks!